



Introduction

RISC-V is an open-source ISA developed by the University of California, Berkeley. The RISC-V ISA takes a modular approach to ISA design and encourages collaboration through practice of openness and transparency. The openness of the ISA plays to its strengths, particularly in the realm of security. Unlike other ISAs, open-source implementations of RISC-V benefit from community-driven security hardening efforts.



OpenBSD is well known for its strong security practices. However, the recent abundance of obscure hardware-level security issues serves as a stark reminder that software systems are only as secure as their underlying hardware. We believe that the combination of security-focused supervisor software (OpenBSD) with an open and transparent hardware platform (RISC-V) will drive forward the development of noticeably more secure platforms for general purpose computing. This project takes the first step towards a full port of the OpenBSD operating system to RISC-V, starting with the kernel.

Architecture

Kernel's Machine-Dependent Components

The project focuses on rewriting the architecture-specific routines and drivers of the machine-dependent layer of the OS kernel to support the RISC-V hardware architecture. The machine-dependent layer encapsulates the hardware details and provides services to the machine-independent layer of OS kernel, which can further provide Application Binary Interface to user programs. The kernel's machine-dependent layer is divided into five categories:


MD Kernel Layer				
Bootstrap	Memory	Trap	Process	Device
Set up paging	<code>pmmap_bootstrap()</code>	Route a trap	PCB	<code>mainbus0</code>
Set up trap entry	<code>pmmap_create()</code>	Timer interrupt	<code>fork()</code>	<code>cpu0</code>
Set up C runtime	<code>pmmap_activate()</code>	External interrupt	Switchframe	<code>com0</code>
Handover to C	<code>pmmap_enter()</code>	Instruction faults	Trapframe	<code>intc0</code>
Set up sigcode	<code>pmmap_remove()</code>	Memory faults	<code>cpu_switch0()</code>	<code>plio0</code>

- Bootstrap: assembly initialization before jumping to C.
- PMAP: manages page table by interacting with MMU.
- Trap: handles interrupts and exceptions.
- Process: forks and schedules processes on HARTs.
- Device: probes and configures devices.

Implementation

Early Bootstrap Subsystem


The early bootstrap subsystem written in assembly mainly sets up a preliminary page table, installs the trap handler, prepares the C runtime and jumps to C routine to continue bootstrap.



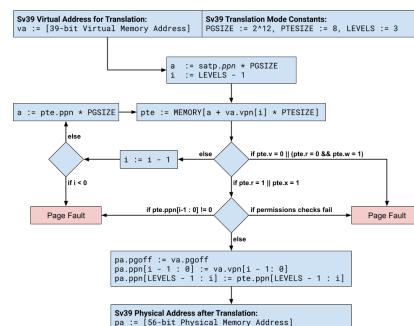
Memory Subsystem

OpenBSD uses the UVM Virtual Memory System to manage system memory resources. UVM operates on RISC-V machines in the Sv39 virtual address translation mode.

Sv39 exposes a page-based virtual memory address scheme with 39 addressable bits. In Sv39 mode, page tables map virtual addresses to physical addresses. Virtual addresses are translated to physical addresses according to the state machine depicted below. The 512 GiB virtual address space is partitioned into regions as shown on the right. The upper 256 GiB is reserved for kernel use and the lower 256 GiB is reserved for user mode processes.




UVM manages the RISC-V MMU via a machine-dependent layer known as the Physical Map (PMAP). In this project, the Physical Map is adapted for the RISC-V MMU in Sv39 mode.



Trap Subsystem

A trap halts the CPU execution and transfers control to a handler for either an interrupt or an exception. Interrupts (software, timer and external interrupts) are managed by a two-level nested interrupt controller architecture as shown below.




Exceptions are synchronous to HART's execution. All supported exceptions can be divided into two categories: instruction faults and memory faults. The OpenBSD kernel utilizes environment call from user-mode to implement system calls. Memory faults are normally handled by a call to `data_abort()`.

Code	Description
0	Instruction address misaligned
1	Instruction access fault
2	Illegal instruction
3	Breakpoint
4	Load address misaligned
5	Load access fault
6	Store / AMO address misaligned
7	Store / AMO access fault
8	Environment call from U-mode
9	Environment call from S-mode
10-11	Reserved
12	Instruction page fault
13	Load page fault
14	Reserved for future standard use
15	Store / AMO page fault
>16	Reserved

Process Subsystem

The process subsystem manages the lifecycle of processes running on the system. A new process is spawned from its parent via a system call to `fork()`. All processes are organized in a tree hierarchy.




CPU scheduler is responsible for switching multiple processes that time-slice a few HARTs. A context switch is performed to save the execution state of the current-running process and restore the execution state of the upcoming process.

Device Subsystem

The device information is presented to kernel via FDT. All devices are probed and configured via the `autocofw` framework. The resultant device hierarchy is shown below.

This project needs to provide the device definition, configuration, and driver files for all MD devices. The `mainbus0` and `simplebus0` are virtual devices introduced for easy device probing and attaching. `timer0` provides periodical interrupts. `intc0` and `plio0` are two nested interrupt controllers. `com0` is the UART device for console output.



Summary/Conclusions

To port OpenBSD to RISC-V ISA, we have implemented assembly initialization for kernel early bootstrap, adapted PMAP to RISC-V MMU in Sv39 mode, implemented routines for process creation and context switch, developed trap handlers for interrupts and exceptions, and developed drivers to probe and attach MD devices.

With supports of RISC-V from the above five subsystems, the OpenBSD kernel can boot up on a single HART and jump to a dummy "Hello-World" userland process.

Key References

- [1] A. Waterman and K. Asanović, "The RISC-V instruction set manual, volume I: Unprivileged ISA," tech. rep., RISC-V Foundation, Dec. 2019.
- [2] A. Waterman and K. Asanović, "The RISC-V instruction set manual, volume II: Privileged architecture," tech. rep., RISC-V Foundation, June 2019.
- [3] H. H. Porter III, "RISC-V: An overview of the instruction set architecture," Jan. 2018.
- [4] C. D. Cranor and G. M. Parulkar, "The UVM virtual memory system," in *Proceedings of the USENIX Annual Technical Conference*, pp. 117–130, June 1999.
- [5] SiFive, "SiFive interrupt cookbook, version 1.2." SiFive, Inc., Feb. 2020.

Acknowledgements

We would like to thank our advisor, Michael Larkin. His expertise with OpenBSD has been essential to the team's success in both the implementation and bring-up of the OpenBSD operating system on RISC-V.